

STAT 583 Lecture 4

Linjun (Leon) Zhang

Department of Statistics
Rutgers University

Feb. 11th

- Midterm is scheduled at 6:50-8:50pm on March 3rd.
- TA: Youmeng Jiang
 - Office hours: 2:00-3:30pm every Monday.

Last time

- Parameter estimation for population mean μ and binomial parameter p
- CI for μ
 - Large n : z-interval
 - Small n : we need to assume the sample is normally distributed
 - when σ is unknown, use Student's t-distribution
 - when σ is known, use normal distribution
- CI for p
 - Large n : z-interval
- CI for $p_1 - p_2$, $\mu_X - \mu_Y$

Recap

- A general formula for CI of mean-based parameters. For a parameter θ , suppose we estimate it by $\hat{\theta}$ and $\hat{\theta}$ is approximated normal. Then a $(1 - \alpha)$ confidence interval for θ is

$$[\hat{\theta} - z_{\alpha/2} \cdot \widehat{\text{sd}}(\hat{\theta}), \quad \hat{\theta} + z_{\alpha/2} \cdot \widehat{\text{sd}}(\hat{\theta})]$$

where $\text{sd}(\hat{\theta})$ is the standard deviation of $\hat{\theta}$, and $\widehat{\text{sd}}(\hat{\theta})$ is the estimate of $\text{sd}(\hat{\theta})$.

- In the setting where the sample size is small and the observations are normally distributed, we replace $z_{\alpha/2}$ with $t_{\alpha/2, n-1}$.

Confidence intervals

- An approximate 95% confidence interval for the binomial parameter p is

$$\left[\hat{p} - 2\sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}, \hat{p} + 2\sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \right]$$

Confidence intervals

- An approximate 95% confidence interval for the binomial parameter p is

$$\left[\hat{p} - 2\sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}, \hat{p} + 2\sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \right]$$

- A conservative 95% confidence interval for $p_1 - p_2$ is

$$\left[\hat{p}_1 - \hat{p}_2 - \sqrt{\frac{1}{n} + \frac{1}{m}}, \hat{p}_1 - \hat{p}_2 + \sqrt{\frac{1}{n} + \frac{1}{m}} \right].$$

Confidence intervals

- A $(1 - \alpha)$ confidence interval for μ is

$$\left[\bar{X} - t_{\alpha/2, n-1} \frac{s}{\sqrt{n}}, \quad \bar{X} + t_{\alpha/2, n-1} \frac{s}{\sqrt{n}} \right].$$

Confidence intervals

- A $(1 - \alpha)$ confidence interval for μ is

$$\left[\bar{X} - t_{\alpha/2, n-1} \frac{s}{\sqrt{n}}, \quad \bar{X} + t_{\alpha/2, n-1} \frac{s}{\sqrt{n}} \right].$$

- A $(1 - \alpha)$ confidence interval for $\mu_X - \mu_Y$ is

$$\left[\bar{X} - \bar{Y} - z_{\alpha/2} \sqrt{\frac{s_X^2}{n} + \frac{s_Y^2}{m}}, \quad \bar{X} - \bar{Y} + z_{\alpha/2} \sqrt{\frac{s_X^2}{n} + \frac{s_Y^2}{m}} \right].$$

Confidence intervals

- A $(1 - \alpha)$ confidence interval for μ is

$$\left[\bar{X} - t_{\alpha/2, n-1} \frac{s}{\sqrt{n}}, \quad \bar{X} + t_{\alpha/2, n-1} \frac{s}{\sqrt{n}} \right].$$

- A $(1 - \alpha)$ confidence interval for $\mu_X - \mu_Y$ is

$$\left[\bar{X} - \bar{Y} - z_{\alpha/2} \sqrt{\frac{s_X^2}{n} + \frac{s_Y^2}{m}}, \quad \bar{X} - \bar{Y} + z_{\alpha/2} \sqrt{\frac{s_X^2}{n} + \frac{s_Y^2}{m}} \right].$$

- A $(1 - \alpha)$ confidence interval for $\mu_X - \mu_Y$ is ($k = \min\{n, m\} - 1$)

$$\left[\bar{X} - \bar{Y} - t_{\alpha/2, k} \sqrt{\frac{s_X^2}{n} + \frac{s_Y^2}{m}}, \quad \bar{X} - \bar{Y} + t_{\alpha/2, k} \sqrt{\frac{s_X^2}{n} + \frac{s_Y^2}{m}} \right].$$

When the variances are the same

Assuming equal variances:

- We combine the information to create what is called a pooled estimate of the variance:

$$s_{pool}^2 = \frac{(n-1)s_X^2 + (m-1)s_Y^2}{n+m-2},$$

$$\text{where } s_X^2 = \frac{(X_1 - \bar{X})^2 + \dots + (X_n - \bar{X})^2}{n-1}, \quad s_Y^2 = \frac{(Y_1 - \bar{Y})^2 + \dots + (Y_m - \bar{Y})^2}{m-1}.$$

When the variances are the same

Assuming equal variances:

- We combine the information to create what is called a pooled estimate of the variance:

$$s_{pool}^2 = \frac{(n-1)s_X^2 + (m-1)s_Y^2}{n+m-2},$$

$$\text{where } s_X^2 = \frac{(X_1 - \bar{X})^2 + \dots + (X_n - \bar{X})^2}{n-1}, \quad s_Y^2 = \frac{(Y_1 - \bar{Y})^2 + \dots + (Y_m - \bar{Y})^2}{m-1}.$$

- A $(1 - \alpha)$ confidence interval for $\mu_X - \mu_Y$ is

$$[\bar{X} - \bar{Y} - t_{\alpha/2, k} \sqrt{\frac{s_p^2}{n} + \frac{s_p^2}{m}}, \bar{X} - \bar{Y} + t_{\alpha/2, k} \sqrt{\frac{s_p^2}{n} + \frac{s_p^2}{m}}],$$

where $k = n + m - 2$.

Paired Sample

- Sometimes it appears we have data from two samples with the further feature that there is a natural “pairing” of the data between the two samples.

Paired Sample

- Sometimes it appears we have data from two samples with the further feature that there is a natural “pairing” of the data between the two samples.
- For example, suppose that the data consists on n brother-sister pairs, with blood pressures X_1, \dots, X_n for the n sisters and blood pressures Y_1, \dots, Y_n for their respective brothers.

CI for the difference between two means

Let's first consider the case where n is small.

- If X_1, X_2, \dots, X_n i.i.d. $\sim N(\mu_X, \sigma_X^2)$, Y_1, Y_2, \dots, Y_n i.i.d. $\sim N(\mu_Y, \sigma_Y^2)$. In the paired samples' case, it's not realistic to assume that $\{X_i\}$ and $\{Y_i\}$ are independent.

CI for the difference between two means

Let's first consider the case where n is small.

- If X_1, X_2, \dots, X_n i.i.d. $\sim N(\mu_X, \sigma_X^2)$, Y_1, Y_2, \dots, Y_n i.i.d. $\sim N(\mu_Y, \sigma_Y^2)$. In the paired samples' case, it's not realistic to assume that $\{X_i\}$ and $\{Y_i\}$ are independent.
- However, it's natural to assume $(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)$ are n independent pairs.

CI for the difference between two means

Let's first consider the case where n is small.

- If X_1, X_2, \dots, X_n i.i.d. $\sim N(\mu_X, \sigma_X^2)$, Y_1, Y_2, \dots, Y_n i.i.d. $\sim N(\mu_Y, \sigma_Y^2)$. In the paired samples' case, it's not realistic to assume that $\{X_i\}$ and $\{Y_i\}$ are independent.
- However, it's natural to assume $(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)$ are n independent pairs.
- We denote $D_i = X_i - Y_i$, then D_i i.i.d. $\sim N(\mu_X - \mu_Y, \sigma_D^2)$

CI for the difference between two means

Let's first consider the case where n is small.

- If X_1, X_2, \dots, X_n i.i.d. $\sim N(\mu_X, \sigma_X^2)$, Y_1, Y_2, \dots, Y_n i.i.d. $\sim N(\mu_Y, \sigma_Y^2)$. In the paired samples' case, it's not realistic to assume that $\{X_i\}$ and $\{Y_i\}$ are independent.
- However, it's natural to assume $(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)$ are n independent pairs.
- We denote $D_i = X_i - Y_i$, then D_i i.i.d. $\sim N(\mu_X - \mu_Y, \sigma_D^2)$
- We estimate $\mu_X - \mu_Y$ by $\bar{D} = \bar{X} - \bar{Y}$

CI for the difference between two means

Let's first consider the case where n is small.

- If X_1, X_2, \dots, X_n i.i.d. $\sim N(\mu_X, \sigma_X^2)$, Y_1, Y_2, \dots, Y_n i.i.d. $\sim N(\mu_Y, \sigma_Y^2)$. In the paired samples' case, it's not realistic to assume that $\{X_i\}$ and $\{Y_i\}$ are independent.
- However, it's natural to assume $(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)$ are n independent pairs.
- We denote $D_i = X_i - Y_i$, then D_i i.i.d. $\sim N(\mu_X - \mu_Y, \sigma_D^2)$
- We estimate $\mu_X - \mu_Y$ by $\bar{D} = \bar{X} - \bar{Y}$
- A $(1 - \alpha)$ confidence interval for $\mu_X - \mu_Y$ is

$$\left[\bar{D} - t_{\alpha/2, n-1} \frac{s_D}{\sqrt{n}}, \quad \bar{D} + t_{\alpha/2, n-1} \frac{s_D}{\sqrt{n}} \right],$$

$$\text{where } s_D = \sqrt{\frac{(D_1 - \bar{D})^2 + \dots + (D_n - \bar{D})^2}{n}}.$$

CI for the difference between two means

Let's first consider the case where n is small.

- If X_1, X_2, \dots, X_n i.i.d. $\sim N(\mu_X, \sigma_X^2)$, Y_1, Y_2, \dots, Y_n i.i.d. $\sim N(\mu_Y, \sigma_Y^2)$. In the paired samples' case, it's not realistic to assume that $\{X_i\}$ and $\{Y_i\}$ are independent.
- However, it's natural to assume $(X_1, Y_1), (X_2, Y_2), \dots, (X_n, Y_n)$ are n independent pairs.
- We denote $D_i = X_i - Y_i$, then D_i i.i.d. $\sim N(\mu_X - \mu_Y, \sigma_D^2)$
- We estimate $\mu_X - \mu_Y$ by $\bar{D} = \bar{X} - \bar{Y}$
- A $(1 - \alpha)$ confidence interval for $\mu_X - \mu_Y$ is

$$\left[\bar{D} - t_{\alpha/2, n-1} \frac{s_D}{\sqrt{n}}, \quad \bar{D} + t_{\alpha/2, n-1} \frac{s_D}{\sqrt{n}} \right],$$

where $s_D = \sqrt{\frac{(D_1 - \bar{D})^2 + \dots + (D_n - \bar{D})^2}{n}}$.

- If n is large, we can simply replace $t_{\alpha/2, n-1}$ with $z_{\alpha/2}$ and remove the normality assumptions.

Example

Suppose that we take the blood pressures of $n = 12$ women and their brothers, and get the following blood pressure reading:

Family	1	2	3	4	5	6	7	8	9	10	11	12
Sister	107	134	111	141	121	118	145	110	164	126	148	132
Brother	110	136	115	140	124	119	148	113	168	129	148	137

Construct the 95% CI for the difference of the mean blood pressure of men and women.

($s_{sister}^2 = 307$, $s_{brother}^2 = 299$, $s_{diff}^2 = 3$, $t_{0.025,11} = 2.306$, $t_{0.05,11} = 1.860$, $t_{0.025,12} = 2.262$, $t_{0.05,12} = 1.833$)

Example

Suppose that we take the blood pressures of $n = 12$ women and their brothers, and get the following blood pressure reading:

Family	1	2	3	4	5	6	7	8	9	10	11	12
Sister	107	134	111	141	121	118	145	110	164	126	148	132
Brother	110	136	115	140	124	119	148	113	168	129	148	137

Construct the 95% CI for the difference of the mean blood pressure of men and women.

($s_{sister}^2 = 307$, $s_{brother}^2 = 299$, $s_{diff}^2 = 3$, $t_{0.025,11} = 2.306$, $t_{0.05,11} = 1.860$, $t_{0.025,12} = 2.262$, $t_{0.05,12} = 1.833$)

Answer

Family	1	2	3	4	5	6	7	8	9	10	11	12
Difference	3	2	4	-1	3	1	3	3	4	3	0	5

• $\bar{D} = 2.5$, $s_D^2 = 3$, $t_{0.025,11} = 2.306$

Example

Suppose that we take the blood pressures of $n = 12$ women and their brothers, and get the following blood pressure reading:

Family	1	2	3	4	5	6	7	8	9	10	11	12
Sister	107	134	111	141	121	118	145	110	164	126	148	132
Brother	110	136	115	140	124	119	148	113	168	129	148	137

Construct the 95% CI for the difference of the mean blood pressure of men and women.

($s_{sister}^2 = 307$, $s_{brother}^2 = 299$, $s_{diff}^2 = 3$, $t_{0.025,11} = 2.306$, $t_{0.05,11} = 1.860$, $t_{0.025,12} = 2.262$, $t_{0.05,12} = 1.833$)

Answer

Family	1	2	3	4	5	6	7	8	9	10	11	12
Difference	3	2	4	-1	3	1	3	3	4	3	0	5

- $\bar{D} = 2.5$, $s_D^2 = 3$, $t_{0.025,11} = 2.306$
- $2.5 \pm 2.306 \cdot \frac{\sqrt{3}}{\sqrt{12}} = [1.35, 3.65]$.

- General confidence intervals construction based on pivotal quantities
- General confidence intervals construction based on maximum likelihood estimators

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**
- A pivotal quantity

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**
- A pivotal quantity
 - is a function of the sample data, the unknown target parameter, and **no other quantities**

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**
- A pivotal quantity
 - is a function of the sample data, the unknown target parameter, and **no other quantities**
 - has a distribution that does not depend on the target parameter

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**
- A pivotal quantity
 - is a function of the sample data, the unknown target parameter, and **no other quantities**
 - has a distribution that does not depend on the target parameter
- Example:

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**
- A pivotal quantity
 - is a function of the sample data, the unknown target parameter, and **no other quantities**
 - has a distribution that does not depend on the target parameter
- Example:
 - We randomly sample an observation from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**
- A pivotal quantity
 - is a function of the sample data, the unknown target parameter, and **no other quantities**
 - has a distribution that does not depend on the target parameter
- Example:
 - We randomly sample an observation from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .
 - If $Y \sim \text{Exp}(\theta)$, then $f_Y(y) = \frac{1}{\theta} e^{-y/\theta}$ for $y \geq 0$.

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**
- A pivotal quantity
 - is a function of the sample data, the unknown target parameter, and **no other quantities**
 - has a distribution that does not depend on the target parameter
- Example:
 - We randomly sample an observation from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .
 - If $Y \sim \text{Exp}(\theta)$, then $f_Y(y) = \frac{1}{\theta} e^{-y/\theta}$ for $y \geq 0$.
 - Let $U = \frac{Y}{\theta}$, we have $f_U(u) = f_Y(u \cdot \theta) \cdot \frac{dy}{du} = \frac{1}{\theta} e^{-u} \cdot \theta = e^{-u}$ for $u > 0$, that is, $U \sim \text{Exp}(1)$

Other confidence intervals

- We have now talked about CI for mean parameters: $\mu, p, \mu_1 - \mu_2, p_1 - p_2$
- For general CIs, a useful method for deriving confidence intervals is to use a **pivotal quantity**
- A pivotal quantity
 - is a function of the sample data, the unknown target parameter, and **no other quantities**
 - has a distribution that does not depend on the target parameter
- Example:
 - We randomly sample an observation from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .
 - If $Y \sim \text{Exp}(\theta)$, then $f_Y(y) = \frac{1}{\theta} e^{-y/\theta}$ for $y \geq 0$.
 - Let $U = \frac{Y}{\theta}$, we have $f_U(u) = f_Y(u \cdot \theta) \cdot \frac{dy}{du} = \frac{1}{\theta} e^{-u} \cdot \theta = e^{-u}$ for $u > 0$, that is, $U \sim \text{Exp}(1)$
 - Thus, we can use $U = \frac{Y}{\theta}$ as a pivotal quantity

Example

- We need to find two numbers a and b such that

$$\mathbb{P}(a \leq U \leq b) = 90\%$$


Example

- We need to find two numbers a and b such that

$$\mathbb{P}(a \leq U \leq b) = 90\%$$

- One way to do this is to choose a and b to satisfy

$$\mathbb{P}(U < a) = \mathbb{P}(U > b) = 5\%$$




Example

- We need to find two numbers a and b such that

$$\mathbb{P}(a \leq U \leq b) = 90\%$$

- One way to do this is to choose a and b to satisfy

$$\mathbb{P}(U < a) = \mathbb{P}(U > b) = 5\%$$



- $1 - e^{-a} = .05$ and $e^{-b} = .05$, equivalently, $a = 0.051$, $b = 2.996$


Example

- We need to find two numbers a and b such that

$$\mathbb{P}(a \leq U \leq b) = 90\%$$

- One way to do this is to choose a and b to satisfy

$$\mathbb{P}(U < a) = \mathbb{P}(U > b) = 5\%$$



- $1 - e^{-a} = .05$ and $e^{-b} = .05$, equivalently, $a = 0.051$, $b = 2.996$
- $0.9 = \mathbb{P}(0.051 \leq \frac{Y}{\theta} \leq 2.996) = \mathbb{P}(\frac{Y}{2.996} \leq \theta \leq \frac{Y}{0.051})$

Example

Example

We randomly sample 10 observations X_1, \dots, X_{10} from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .

Example

Example

We randomly sample 10 observations X_1, \dots, X_{10} from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .

Answer

- Hint: If $X_i \sim \text{Exp}(\theta_i)$ then $\min\{X_1, \dots, X_n\} \sim \text{Exp}\left(\frac{1}{\frac{1}{\theta_1} + \dots + \frac{1}{\theta_n}}\right)$.

Example

Example

We randomly sample 10 observations X_1, \dots, X_{10} from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .

Answer

- Hint: If $X_i \sim \text{Exp}(\theta_i)$ then $\min\{X_1, \dots, X_n\} \sim \text{Exp}\left(\frac{1}{\frac{1}{\theta_1} + \dots + \frac{1}{\theta_n}}\right)$.
- $\frac{X_1}{\theta}, \dots, \frac{X_{10}}{\theta} \stackrel{i.i.d.}{\sim} \text{Exp}(1)$

Example

Example

We randomly sample 10 observations X_1, \dots, X_{10} from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .

Answer

- Hint: If $X_i \sim \text{Exp}(\theta_i)$ then $\min\{X_1, \dots, X_n\} \sim \text{Exp}(\frac{1}{\frac{1}{\theta_1} + \dots + \frac{1}{\theta_n}})$.
- $\frac{X_1}{\theta}, \dots, \frac{X_{10}}{\theta} \stackrel{i.i.d.}{\sim} \text{Exp}(1)$
- $U = \min\{\frac{X_1}{\theta}, \dots, \frac{X_{10}}{\theta}\} \sim \text{Exp}(\frac{1}{10})$

Example

Example

We randomly sample 10 observations X_1, \dots, X_{10} from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .

Answer

- Hint: If $X_i \sim \text{Exp}(\theta_i)$ then $\min\{X_1, \dots, X_n\} \sim \text{Exp}(\frac{1}{\frac{1}{\theta_1} + \dots + \frac{1}{\theta_n}})$.
- $\frac{X_1}{\theta}, \dots, \frac{X_{10}}{\theta} \stackrel{i.i.d.}{\sim} \text{Exp}(1)$
- $U = \min\{\frac{X_1}{\theta}, \dots, \frac{X_{10}}{\theta}\} \sim \text{Exp}(\frac{1}{10})$
- Find a, b such that $\mathbb{P}(U < a) = 1 - e^{-10a} = .05$ and $\mathbb{P}(U > b) = e^{-10b} = .05$, equivalently, $a = 0.005$, $b = 0.300$

Example

Example

We randomly sample 10 observations X_1, \dots, X_{10} from an exponential distribution with unknown mean θ . Find a formula for a 90% CI for θ .

Answer

- Hint: If $X_i \sim \text{Exp}(\theta_i)$ then $\min\{X_1, \dots, X_n\} \sim \text{Exp}(\frac{1}{\frac{1}{\theta_1} + \dots + \frac{1}{\theta_n}})$.
- $\frac{X_1}{\theta}, \dots, \frac{X_{10}}{\theta} \stackrel{i.i.d.}{\sim} \text{Exp}(1)$
- $U = \min\{\frac{X_1}{\theta}, \dots, \frac{X_{10}}{\theta}\} \sim \text{Exp}(\frac{1}{10})$
- Find a, b such that $\mathbb{P}(U < a) = 1 - e^{-10a} = .05$ and $\mathbb{P}(U > b) = e^{-10b} = .05$, equivalently, $a = 0.005$, $b = 0.300$
- $0.9 = \mathbb{P}(0.005 < U < 0.300) = \mathbb{P}(0.005 < \frac{\min\{X_1, \dots, X_{10}\}}{\theta} < 0.300) = \mathbb{P}(\frac{\min\{X_1, \dots, X_{10}\}}{0.300} < \theta < \frac{\min\{X_1, \dots, X_{10}\}}{0.005})$

Example

Example

Suppose that we take a sample Y from a uniform distribution defined on the interval $[0, \theta]$, where θ is unknown. Find a 95% lower confidence bound for θ , that is, find $L(Y)$, such that $\mathbb{P}(\theta \geq L(Y)) = 0.95$.

Example

Example

Suppose that we take a sample Y from a uniform distribution defined on the interval $[0, \theta]$, where θ is unknown. Find a 95% lower confidence bound for θ , that is, find $L(Y)$, such that $\mathbb{P}(\theta \geq L(Y)) = 0.95$.

Answer

- Consider $U = \frac{Y}{\theta}$

Example

Example

Suppose that we take a sample Y from a uniform distribution defined on the interval $[0, \theta]$, where θ is unknown. Find a 95% lower confidence bound for θ , that is, find $L(Y)$, such that $\mathbb{P}(\theta \geq L(Y)) = 0.95$.

Answer

- Consider $U = \frac{Y}{\theta}$
- $U \sim U(0, 1)$

Example

Example

Suppose that we take a sample Y from a uniform distribution defined on the interval $[0, \theta]$, where θ is unknown. Find a 95% lower confidence bound for θ , that is, find $L(Y)$, such that $\mathbb{P}(\theta \geq L(Y)) = 0.95$.

Answer

- Consider $U = \frac{Y}{\theta}$
- $U \sim U(0, 1)$
- $\mathbb{P}(U < 0.95) = 0.95$

Example

Example

Suppose that we take a sample Y from a uniform distribution defined on the interval $[0, \theta]$, where θ is unknown. Find a 95% lower confidence bound for θ , that is, find $L(Y)$, such that $\mathbb{P}(\theta \geq L(Y)) = 0.95$.

Answer

- Consider $U = \frac{Y}{\theta}$
- $U \sim U(0, 1)$
- $\mathbb{P}(U < 0.95) = 0.95$
- $0.95 = \mathbb{P}(U < 0.95) = \mathbb{P}\left(\frac{Y}{\theta} < 0.95\right) = \mathbb{P}\left(\frac{Y}{0.95} \leq \theta\right)$

Example

Example

Let Y_1, \dots, Y_n denote a random sample of size n from $U(0, \theta)$. Let $Y_{(n)} = \max\{Y_1, Y_2, \dots, Y_n\}$. Find a 95% lower confidence bound for θ .

Example

Example

Let Y_1, \dots, Y_n denote a random sample of size n from $U(0, \theta)$. Let $Y_{(n)} = \max\{Y_1, Y_2, \dots, Y_n\}$. Find a 95% lower confidence bound for θ .

Answer

- Consider $U = \frac{Y_{(n)}}{\theta}$

Example

Example

Let Y_1, \dots, Y_n denote a random sample of size n from $U(0, \theta)$. Let $Y_{(n)} = \max\{Y_1, Y_2, \dots, Y_n\}$. Find a 95% lower confidence bound for θ .

Answer

- Consider $U = \frac{Y_{(n)}}{\theta}$
- U has the cdf
$$F(u) = \mathbb{P}\left(\frac{\max\{Y_1, Y_2, \dots, Y_n\}}{\theta} \leq u\right) = \mathbb{P}\left(\frac{Y_1}{\theta} \leq u, \dots, \frac{Y_n}{\theta} \leq u\right) = \mathbb{P}\left(\frac{Y_1}{\theta} \leq u\right)^n$$

Example

Example

Let Y_1, \dots, Y_n denote a random sample of size n from $U(0, \theta)$. Let $Y_{(n)} = \max\{Y_1, Y_2, \dots, Y_n\}$. Find a 95% lower confidence bound for θ .

Answer

- Consider $U = \frac{Y_{(n)}}{\theta}$
- U has the cdf
$$F(u) = \mathbb{P}\left(\frac{\max\{Y_1, Y_2, \dots, Y_n\}}{\theta} \leq u\right) = \mathbb{P}\left(\frac{Y_1}{\theta} \leq u, \dots, \frac{Y_n}{\theta} \leq u\right) = \mathbb{P}\left(\frac{Y_1}{\theta} \leq u\right)^n$$
- $F(u) = u^n$ for $u \in [0, 1]$.

Example

Example

Let Y_1, \dots, Y_n denote a random sample of size n from $U(0, \theta)$. Let $Y_{(n)} = \max\{Y_1, Y_2, \dots, Y_n\}$. Find a 95% lower confidence bound for θ .

Answer

- Consider $U = \frac{Y_{(n)}}{\theta}$
- U has the cdf
$$F(u) = \mathbb{P}\left(\frac{\max\{Y_1, Y_2, \dots, Y_n\}}{\theta} \leq u\right) = \mathbb{P}\left(\frac{Y_1}{\theta} \leq u, \dots, \frac{Y_n}{\theta} \leq u\right) = \mathbb{P}\left(\frac{Y_1}{\theta} \leq u\right)^n$$
- $F(u) = u^n$ for $u \in [0, 1]$.
- $0.95 = \mathbb{P}(U < (0.95)^{1/n}) = \mathbb{P}\left(\frac{Y_{(n)}}{\theta} < (0.95)^{1/n}\right) = \mathbb{P}\left(\frac{Y_{(n)}}{(0.95)^{1/n}} \leq \theta\right)$

Example

Example

If we have $n = 10$ samples i.i.d. from $U(0, \theta)$, with the max $Y_{\max} = 5.7$. We are 95% confident that θ is at least:.

Example

Example

If we have $n = 10$ samples i.i.d. from $U(0, \theta)$, with the max $Y_{\max} = 5.7$. We are 95% confident that θ is at least:.

Answer

$$\bullet 0.95 = \mathbb{P}(U < (0.95)^{1/n}) = \mathbb{P}\left(\frac{Y_{(n)}}{\theta} < (0.95)^{1/n}\right) = \mathbb{P}\left(\frac{Y_{(n)}}{(0.95)^{1/n}} \leq \theta\right)$$

Example

Example

If we have $n = 10$ samples i.i.d. from $U(0, \theta)$, with the max $Y_{\max} = 5.7$. We are 95% confident that θ is at least:.

Answer

- $0.95 = \mathbb{P}(U < (0.95)^{1/n}) = \mathbb{P}(\frac{Y_{(n)}}{\theta} < (0.95)^{1/n}) = \mathbb{P}(\frac{Y_{(n)}}{(0.95)^{1/n}} \leq \theta)$
- $\frac{Y_{(n)}}{(0.95)^{1/n}} = \frac{5.7}{(0.95)^{1/10}} = 5.73$

Confidence interval for the variance

- Suppose $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$ with μ and σ^2 unknown. We seek a $100(1 - \alpha)\%$ CI for σ^2

Confidence interval for the variance

- Suppose $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$ with μ and σ^2 unknown. We seek a $100(1 - \alpha)\%$ CI for σ^2
- Recall that, for $S^2 = \frac{1}{n-1} \sum_{i=1}^n (Y_i - \bar{Y})^2$, we have the pivotal quantity

$$\frac{(n-1)S^2}{\sigma^2} \sim \chi_{n-1}^2.$$

Confidence interval for the variance

- Suppose $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$ with μ and σ^2 unknown. We seek a $100(1 - \alpha)\%$ CI for σ^2
- Recall that, for $S^2 = \frac{1}{n-1} \sum_{i=1}^n (Y_i - \bar{Y})^2$, we have the pivotal quantity

$$\frac{(n-1)S^2}{\sigma^2} \sim \chi_{n-1}^2.$$

- As a result,

$$\mathbb{P}(\chi_{1-(\alpha/2), n-1}^2 \leq \frac{(n-1)S^2}{\sigma^2} \leq \chi_{\alpha/2, n-1}^2) = 1 - \alpha$$

Confidence interval for the variance

- Suppose $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$ with μ and σ^2 unknown. We seek a $100(1 - \alpha)\%$ CI for σ^2
- Recall that, for $S^2 = \frac{1}{n-1} \sum_{i=1}^n (Y_i - \bar{Y})^2$, we have the pivotal quantity

$$\frac{(n-1)S^2}{\sigma^2} \sim \chi_{n-1}^2.$$

- As a result,

$$\mathbb{P}(\chi_{1-(\alpha/2), n-1}^2 \leq \frac{(n-1)S^2}{\sigma^2} \leq \chi_{\alpha/2, n-1}^2) = 1 - \alpha$$

- A $100(1 - \alpha)\%$ Confidence Interval for σ^2 :

$$\mathbb{P}\left(\frac{(n-1)S^2}{\chi_{\alpha/2, n-1}^2} \leq \sigma^2 \leq \frac{(n-1)S^2}{\chi_{1-(\alpha/2), n-1}^2}\right) = 1 - \alpha.$$

Example

Example

Suppose the maturation times for a flower species are $N(\mu, \sigma^2)$. If a random sample of $n = 13$ seeds yielded $s^2 = 10.7$, then what is a 90% CI for σ^2 ?

($\chi^2_{0.05,12} = 21.03$, $\chi^2_{0.95,12} = 5.23$)

Example

Example

Suppose the maturation times for a flower species are $N(\mu, \sigma^2)$. If a random sample of $n = 13$ seeds yielded $s^2 = 10.7$, then what is a 90% CI for σ^2 ?
($\chi_{0.05,12}^2 = 21.03$, $\chi_{0.95,12}^2 = 5.23$)

Answer

$$\mathbb{P}\left(\frac{(n-1)S^2}{\chi_{\alpha/2, n-1}^2} \leq \sigma^2 \leq \frac{(n-1)S^2}{\chi_{1-(\alpha/2), n-1}^2}\right) = 1 - \alpha.$$

Example

Example

Suppose the maturation times for a flower species are $N(\mu, \sigma^2)$. If a random sample of $n = 13$ seeds yielded $s^2 = 10.7$, then what is a 90% CI for σ^2 ?
($\chi_{0.05,12}^2 = 21.03$, $\chi_{0.95,12}^2 = 5.23$)

Answer

$$\mathbb{P}\left(\frac{(n-1)S^2}{\chi_{\alpha/2,n-1}^2} \leq \sigma^2 \leq \frac{(n-1)S^2}{\chi_{1-(\alpha/2),n-1}^2}\right) = 1 - \alpha.$$

$$\frac{(n-1)S^2}{\chi_{\alpha/2,n-1}^2} = \frac{12 \times 10.7}{21.03} = 6.11$$

Example

Example

Suppose the maturation times for a flower species are $N(\mu, \sigma^2)$. If a random sample of $n = 13$ seeds yielded $s^2 = 10.7$, then what is a 90% CI for σ^2 ?
($\chi_{0.05,12}^2 = 21.03$, $\chi_{0.95,12}^2 = 5.23$)

Answer

- $$\mathbb{P}\left(\frac{(n-1)S^2}{\chi_{\alpha/2,n-1}^2} \leq \sigma^2 \leq \frac{(n-1)S^2}{\chi_{1-(\alpha/2),n-1}^2}\right) = 1 - \alpha.$$
- $$\frac{(n-1)S^2}{\chi_{\alpha/2,n-1}^2} = \frac{12 \times 10.7}{21.03} = 6.11$$
- $$\frac{(n-1)S^2}{\chi_{1-\alpha/2,n-1}^2} = \frac{12 \times 10.7}{5.23} = 24.55$$

Example

Example

Suppose the maturation times for a flower species are $N(\mu, \sigma^2)$. If a random sample of $n = 13$ seeds yielded $s^2 = 10.7$, then what is a 90% CI for σ^2 ?

($\chi_{0.05,12}^2 = 21.03$, $\chi_{0.95,12}^2 = 5.23$)

Answer



$$\mathbb{P}\left(\frac{(n-1)S^2}{\chi_{\alpha/2,n-1}^2} \leq \sigma^2 \leq \frac{(n-1)S^2}{\chi_{1-(\alpha/2),n-1}^2}\right) = 1 - \alpha.$$

- $\frac{(n-1)S^2}{\chi_{\alpha/2,n-1}^2} = \frac{12 \times 10.7}{21.03} = 6.11$
- $\frac{(n-1)S^2}{\chi_{1-\alpha/2,n-1}^2} = \frac{12 \times 10.7}{5.23} = 24.55$
- The 90% CI is [6.11, 24.55].

Maximum likelihood estimator

- Another general and important method to construct point estimation and confidence intervals are **maximum likelihood estimator (MLE)**

Likelihood function

- Suppose Y_1, \dots, Y_n are an i.i.d. random sample from pdf $f_Y(y; \theta)$

Likelihood function

- Suppose Y_1, \dots, Y_n are an i.i.d. random sample from pdf $f_Y(y; \theta)$
- Then we call

$$f(y_1, y_2, \dots, y_n | \theta)$$

the **joint density function**. It is a function of y_1, \dots, y_n for a fixed (unknown) value of θ .

Likelihood function

- Suppose Y_1, \dots, Y_n are an i.i.d. random sample from pdf $f_Y(y; \theta)$
- Then we call

$$f(y_1, y_2, \dots, y_n | \theta)$$

the **joint density function**. It is a function of y_1, \dots, y_n for a fixed (unknown) value of θ .

- This function gives the probability or likelihood of observing the event $\{(Y_1 = y_1, Y_2 = y_2, \dots, Y_n = y_n)\}$ when the value of the parameter is θ .

Likelihood function

- Suppose Y_1, \dots, Y_n are an i.i.d. random sample from pdf $f_Y(y; \theta)$
- Then we call

$$f(y_1, y_2, \dots, y_n | \theta)$$

the **joint density function**. It is a function of y_1, \dots, y_n for a fixed (unknown) value of θ .

- This function gives the probability or likelihood of observing the event $\{(Y_1 = y_1, Y_2 = y_2, \dots, Y_n = y_n)\}$ when the value of the parameter is θ .
- By the i.i.d. assumption, we have

$$f(y_1, y_2, \dots, y_n | \theta) = \prod_{i=1}^n f(y_i | \theta)$$

Likelihood function

- We could also view $f(y_1, y_2, \dots, y_n | \theta)$ as a function of θ for a given data set y_1, \dots, y_n . In that case we write

$$L(\theta \mid y_1, \dots, y_n) := f(y_1, y_2, \dots, y_n | \theta).$$

Likelihood function

- We could also view $f(y_1, y_2, \dots, y_n | \theta)$ as a function of θ for a given data set y_1, \dots, y_n . In that case we write

$$L(\theta \mid y_1, \dots, y_n) := f(y_1, y_2, \dots, y_n | \theta).$$

- We call $L(\theta \mid y_1, \dots, y_n)$ the **likelihood function**.

Likelihood function

- We could also view $f(y_1, y_2, \dots, y_n | \theta)$ as a function of θ for a given data set y_1, \dots, y_n . In that case we write

$$L(\theta | y_1, \dots, y_n) := f(y_1, y_2, \dots, y_n | \theta).$$

- We call $L(\theta | y_1, \dots, y_n)$ the **likelihood function**.
- Note: the joint density and likelihood function are the same function, but the first is treated as a function of the data, and the second as a function of the parameter θ .

Maximum likelihood estimator

- Loosely speaking, the value of θ that maximizes $L(\theta \mid y_1, \dots, y_n)$ is the parameter value that is “most likely” to have produced the data we observe.

Maximum likelihood estimator

- Loosely speaking, the value of θ that maximizes $L(\theta \mid y_1, \dots, y_n)$ is the parameter value that is “most likely” to have produced the data we observe.
- The **maximum likelihood estimator (MLE)** of the parameter θ is

$$\hat{\theta} = \arg \max L(\theta \mid y_1, \dots, y_n)$$

Maximum likelihood estimator

- Loosely speaking, the value of θ that maximizes $L(\theta \mid y_1, \dots, y_n)$ is the parameter value that is “most likely” to have produced the data we observe.
- The **maximum likelihood estimator (MLE)** of the parameter θ is

$$\hat{\theta} = \arg \max L(\theta \mid y_1, \dots, y_n)$$

- Note: θ can be taken as a vector $\theta = (\theta_1, \dots, \theta_k)$

Maximum likelihood estimator

- Loosely speaking, the value of θ that maximizes $L(\theta \mid y_1, \dots, y_n)$ is the parameter value that is “most likely” to have produced the data we observe.
- The **maximum likelihood estimator (MLE)** of the parameter θ is

$$\hat{\theta} = \arg \max L(\theta \mid y_1, \dots, y_n)$$

- Note: θ can be taken as a vector $\theta = (\theta_1, \dots, \theta_k)$
- Note 2: In many cases, it is easier to maximize $\log L(\theta)$ than to maximize $L(\theta)$ itself, since log is a strictly increasing function. We call $\log L(\theta)$ the **log likelihood**.

Maximum likelihood estimator

Often, the MLE is found by

- 1 Writing out the (log) likelihood as a function of the parameter (say, θ)
- 2 Taking the derivative with respect to θ
- 3 Setting the derivative equal to 0 and solving for $\hat{\theta}$
- 4 Checking that the second derivative is **negative** at $\hat{\theta}$ to ensure the solution is a maximum.

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p
- Solution

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p
- Solution
 - $f(y_i \mid p) = p^{y_i}(1 - p)^{1-y_i}$.

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p
- Solution
 - $f(y_i | p) = p^{y_i}(1 - p)^{1-y_i}$.
 - $L(p) = L(y_1, y_2, \dots, y_n | p) = p^y(1 - p)^{n-y}$, where $y = \sum_{i=1}^n y_i$.

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p
- Solution
 - $f(y_i | p) = p^{y_i}(1 - p)^{1-y_i}$.
 - $L(p) = L(y_1, y_2, \dots, y_n | p) = p^y(1 - p)^{n-y}$, where $y = \sum_{i=1}^n y_i$.
 - $\log L(p) = y \log p + (n - y) \log(1 - p)$

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p

- Solution

- $f(y_i | p) = p^{y_i}(1-p)^{1-y_i}$.
- $L(p) = L(y_1, y_2, \dots, y_n | p) = p^y(1-p)^{n-y}$, where $y = \sum_{i=1}^n y_i$.
- $\log L(p) = y \log p + (n-y) \log(1-p)$
- $\frac{d(\log L(p))}{dp} = y\left(\frac{1}{p}\right) + (n-y)\frac{-1}{1-p}$

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p

- Solution

- $f(y_i | p) = p^{y_i}(1 - p)^{1 - y_i}$.
- $L(p) = L(y_1, y_2, \dots, y_n | p) = p^y(1 - p)^{n - y}$, where $y = \sum_{i=1}^n y_i$.
- $\log L(p) = y \log p + (n - y) \log(1 - p)$
- $\frac{d(\log L(p))}{dp} = y\left(\frac{1}{p}\right) + (n - y)\frac{-1}{n - p}$
- We need to solve for \hat{p} such that

$$\frac{y}{\hat{p}} = \frac{n - y}{1 - \hat{p}}$$

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p

- Solution

- $f(y_i | p) = p^{y_i}(1 - p)^{1 - y_i}$.
- $L(p) = L(y_1, y_2, \dots, y_n | p) = p^y(1 - p)^{n - y}$, where $y = \sum_{i=1}^n y_i$.
- $\log L(p) = y \log p + (n - y) \log(1 - p)$
- $\frac{d(\log L(p))}{dp} = y\left(\frac{1}{p}\right) + (n - y)\frac{-1}{n - p}$
- We need to solve for \hat{p} such that

$$\frac{y}{\hat{p}} = \frac{n - y}{1 - \hat{p}}$$

- We then have $\hat{p} = \frac{y}{n}$

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of p

- Solution

- $f(y_i | p) = p^{y_i}(1 - p)^{1-y_i}$.
- $L(p) = L(y_1, y_2, \dots, y_n | p) = p^y(1 - p)^{n-y}$, where $y = \sum_{i=1}^n y_i$.
- $\log L(p) = y \log p + (n - y) \log(1 - p)$
- $\frac{d(\log L(p))}{dp} = y\left(\frac{1}{p}\right) + (n - y)\frac{-1}{n-p}$
- We need to solve for \hat{p} such that

$$\frac{y}{\hat{p}} = \frac{n - y}{1 - \hat{p}}$$

- We then have $\hat{p} = \frac{y}{n}$
- Hence, the MLE of p is actually the intuitive estimator for p that we used before

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} U(0, \theta)$. Find the MLE of θ

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} U(0, \theta)$. Find the MLE of θ
- Solution

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} U(0, \theta)$. Find the MLE of θ
- Solution
 - $f(y_i | \theta) = \frac{1}{\theta} I_{[0, \theta]}(y_i)$, , where $I_{[a, b]}(\cdot)$ is an indicator function such that $I_{a, b}(x) = 1$ if $x \in [a, b]$, and $= 0$ otherwise

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} U(0, \theta)$. Find the MLE of θ
- Solution
 - $f(y_i | \theta) = \frac{1}{\theta} I_{[0, \theta]}(y_i)$, where $I_{[a, b]}(\cdot)$ is an indicator function such that $I_{a, b}(x) = 1$ if $x \in [a, b]$, and $= 0$ otherwise
 - $L(\theta) = L(y_1, y_2, \dots, y_n | \theta) = \prod_{i=1}^n I_{[0, \theta]}(y_i) = \frac{1}{\theta^n}$ if $y_1, \dots, y_n \in [0, \theta]$, and 0 otherwise

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} U(0, \theta)$. Find the MLE of θ
- Solution
 - $f(y_i | \theta) = \frac{1}{\theta} I_{[0, \theta]}(y_i)$, , where $I_{[a, b]}(\cdot)$ is an indicator function such that $I_{a, b}(x) = 1$ if $x \in [a, b]$, and $= 0$ otherwise
 - $L(\theta) = L(y_1, y_2, \dots, y_n | \theta) = \prod_{i=1}^n I_{[0, \theta]}(y_i) = \frac{1}{\theta^n}$ if $y_1, \dots, y_n \in [0, \theta]$, and 0 otherwise
 - $L(\theta)$ is maximized if $\hat{\theta} = \max_i y_i$

Example of maximum likelihood estimator

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} U(0, \theta)$. Find the MLE of θ
- Solution
 - $f(y_i | \theta) = \frac{1}{\theta} I_{[0, \theta]}(y_i)$, , where $I_{[a, b]}(\cdot)$ is an indicator function such that $I_{a, b}(x) = 1$ if $x \in [a, b]$, and $= 0$ otherwise
 - $L(\theta) = L(y_1, y_2, \dots, y_n | \theta) = \prod_{i=1}^n I_{[0, \theta]}(y_i) = \frac{1}{\theta^n}$ if $y_1, \dots, y_n \in [0, \theta]$, and 0 otherwise
 - $L(\theta)$ is maximized if $\hat{\theta} = \max_i y_i$
 - The MLE is $\max_i y_i$

Example of maximum likelihood estimator

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Exp}(\theta)$. Find the MLE of θ .

Solution

- $f(y_i | \theta) = \frac{1}{\theta} e^{-y_i/\theta}$

Example of maximum likelihood estimator

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Exp}(\theta)$. Find the MLE of θ .

Solution

- $f(y_i | \theta) = \frac{1}{\theta} e^{-y_i/\theta}$
- $L(\theta) = L(y_1, y_2, \dots, y_n | \theta) = \theta^{-n} e^{-(y_1 + \dots + y_n)/\theta}$

Example of maximum likelihood estimator

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Exp}(\theta)$. Find the MLE of θ .

Solution

- $f(y_i | \theta) = \frac{1}{\theta} e^{-y_i/\theta}$
- $L(\theta) = L(y_1, y_2, \dots, y_n | \theta) = \theta^{-n} e^{-(y_1 + \dots + y_n)/\theta}$
- $\log L(\theta) = -n \log \theta - (y_1 + \dots + y_n)/\theta$

Example of maximum likelihood estimator

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Exp}(\theta)$. Find the MLE of θ .

Solution

- $f(y_i | \theta) = \frac{1}{\theta} e^{-y_i/\theta}$
- $L(\theta) = L(y_1, y_2, \dots, y_n | \theta) = \theta^{-n} e^{-(y_1 + \dots + y_n)/\theta}$
- $\log L(\theta) = -n \log \theta - (y_1 + \dots + y_n)/\theta$
- $\frac{d(\log L(\theta))}{d\theta} = -\frac{n}{\theta} + \frac{y_1 + \dots + y_n}{\theta^2}$

Example of maximum likelihood estimator

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Exp}(\theta)$. Find the MLE of θ .

Solution

- $f(y_i | \theta) = \frac{1}{\theta} e^{-y_i/\theta}$
- $L(\theta) = L(y_1, y_2, \dots, y_n | \theta) = \theta^{-n} e^{-(y_1 + \dots + y_n)/\theta}$
- $\log L(\theta) = -n \log \theta - (y_1 + \dots + y_n)/\theta$
- $\frac{d(\log L(\theta))}{d\theta} = -\frac{n}{\theta} + \frac{y_1 + \dots + y_n}{\theta^2}$
- We need to solve for $\hat{\theta}$ such that $\frac{n}{\hat{\theta}} = \frac{y_1 + \dots + y_n}{\hat{\theta}^2}$

Example of maximum likelihood estimator

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Exp}(\theta)$. Find the MLE of θ .

Solution

- $f(y_i | \theta) = \frac{1}{\theta} e^{-y_i/\theta}$
- $L(\theta) = L(y_1, y_2, \dots, y_n | \theta) = \theta^{-n} e^{-(y_1 + \dots + y_n)/\theta}$
- $\log L(\theta) = -n \log \theta - (y_1 + \dots + y_n)/\theta$
- $\frac{d(\log L(\theta))}{d\theta} = -\frac{n}{\theta} + \frac{y_1 + \dots + y_n}{\theta^2}$
- We need to solve for $\hat{\theta}$ such that $\frac{n}{\hat{\theta}} = \frac{y_1 + \dots + y_n}{\hat{\theta}^2}$
- We then have $\hat{\theta} = \frac{y_1 + \dots + y_n}{n}$.

Example of MLE with multiple parameters

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$. Find the MLE of μ and σ .

Example of MLE with multiple parameters

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$. Find the MLE of μ and σ .
- $f(y_i \mid \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left[-\frac{(y_i - \mu)^2}{2\sigma^2}\right]$

Example of MLE with multiple parameters

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$. Find the MLE of μ and σ .
- $f(y_i | \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left[-\frac{(y_i - \mu)^2}{2\sigma^2}\right]$
- $L(\lambda) = L(y_1, y_2, \dots, y_n | \lambda) = \left(\frac{1}{2\pi\sigma^2}\right)^{n/2} \exp\left[-\frac{\sum_{i=1}^n (y_i - \mu)^2}{2\sigma^2}\right]$

Example of MLE with multiple parameters

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$. Find the MLE of μ and σ .
- $f(y_i | \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp[-\frac{(y_i - \mu)^2}{2\sigma^2}]$
- $L(\lambda) = L(y_1, y_2, \dots, y_n | \lambda) = (\frac{1}{2\pi\sigma^2})^{n/2} \exp[-\frac{\sum_{i=1}^n (y_i - \mu)^2}{2\sigma^2}]$
- $\log L(\lambda) = -\frac{n}{2} \log \sigma^2 - \frac{n}{2} \log 2\pi - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mu)^2$.

Example of MLE with multiple parameters

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$. Find the MLE of μ and σ .
- $f(y_i | \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp[-\frac{(y_i - \mu)^2}{2\sigma^2}]$
- $L(\lambda) = L(y_1, y_2, \dots, y_n | \lambda) = (\frac{1}{2\pi\sigma^2})^{n/2} \exp[-\frac{\sum_{i=1}^n (y_i - \mu)^2}{2\sigma^2}]$
- $\log L(\lambda) = -\frac{n}{2} \log \sigma^2 - \frac{n}{2} \log 2\pi - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mu)^2$.
- Taking derivatives:

$$\frac{\partial(\log L(\mu, \sigma^2))}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^n (y_i - \mu)$$

$$\frac{\partial(\log L(\mu, \sigma^2))}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (y_i - \mu)^2$$

Example of MLE with multiple parameters

- Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} N(\mu, \sigma^2)$. Find the MLE of μ and σ .
- $f(y_i | \mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp[-\frac{(y_i - \mu)^2}{2\sigma^2}]$
- $L(\lambda) = L(y_1, y_2, \dots, y_n | \lambda) = (\frac{1}{2\pi\sigma^2})^{n/2} \exp[-\frac{\sum_{i=1}^n (y_i - \mu)^2}{2\sigma^2}]$
- $\log L(\lambda) = -\frac{n}{2} \log \sigma^2 - \frac{n}{2} \log 2\pi - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - \mu)^2$.
- Taking derivatives:

$$\frac{\partial(\log L(\mu, \sigma^2))}{\partial \mu} = \frac{1}{\sigma^2} \sum_{i=1}^n (y_i - \mu)$$

$$\frac{\partial(\log L(\mu, \sigma^2))}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n (y_i - \mu)^2$$

- Setting these derivatives equal to zero, we get $\mu = \frac{1}{n} \sum_{i=1}^n y_i$ and $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \mu)^2$.

Properties of MLEs

- We are often interested in estimating a function of a parameter $g(\theta)$
- The **invariance property** of MLEs states that if $\hat{\theta}$ is a MLE of θ , and $g(\cdot)$ is any function, then $g(\hat{\theta})$ is a MLE of $g(\theta)$.

Example

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of $\text{Var}(\sum_{i=1}^n Y_i)$.

Solution

- $\text{Var}(\sum_{i=1}^n y_i) = n\text{Var}(y_1) = np(1 - p)$

Example

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of $\text{Var}(\sum_{i=1}^n Y_i)$.

Solution

- $\text{Var}(\sum_{i=1}^n y_i) = n\text{Var}(y_1) = np(1 - p)$
- The MLE of p is $\hat{p} = \frac{1}{n} \sum_{i=1}^n Y_i$

Example

Problem

Suppose that $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} \text{Ber}(p)$. Find the MLE of $\text{Var}(\sum_{i=1}^n Y_i)$.

Solution

- $\text{Var}(\sum_{i=1}^n y_i) = n\text{Var}(y_1) = np(1 - p)$
- The MLE of p is $\hat{p} = \frac{1}{n} \sum_{i=1}^n Y_i$
- The MLE of $np(1 - p)$ is $n\hat{p}(1 - \hat{p}) = \sum_{i=1}^n Y_i(1 - \frac{1}{n} \sum_{i=1}^n Y_i)$

Asymptotic properties of estimators

- The asymptotic behavior of an estimator refers to its behavior as the sample size n increases towards ∞

Asymptotic properties of estimators

- The asymptotic behavior of an estimator refers to its behavior as the sample size n increases towards ∞
- Since an estimator $\hat{\theta}$ usually depends on n , we may denote a sequence of such estimators by $\hat{\theta}_n$

Asymptotic properties of estimators

- The asymptotic behavior of an estimator refers to its behavior as the sample size n increases towards ∞
- Since an estimator $\hat{\theta}$ usually depends on n , we may denote a sequence of such estimators by $\hat{\theta}_n$
- Intuitively, we would like our estimator to “get closer” to the target parameter as $n \rightarrow \infty$

Asymptotic properties of estimators

- The asymptotic behavior of an estimator refers to its behavior as the sample size n increases towards ∞
- Since an estimator $\hat{\theta}$ usually depends on n , we may denote a sequence of such estimators by $\hat{\theta}_n$
- Intuitively, we would like our estimator to “get closer” to the target parameter as $n \rightarrow \infty$
- Definition: An estimator $\hat{\theta}_n$ is a **consistent** estimator of θ if, for any $\epsilon > 0$,

$$\lim_{n \rightarrow \infty} \mathbb{P}(|\hat{\theta}_n - \theta| \leq \epsilon) = 1,$$

also denote as $\hat{\theta}_n \xrightarrow{P} \theta$ (converge in probability).

Asymptotic properties of MLE

Theorem

If $Y_1, \dots, Y_n \stackrel{i.i.d.}{\sim} f_Y(y; \theta)$ and $\hat{\theta}$ is the MLE of θ , then assuming certain regularity conditions:

$$\hat{\theta}_n \xrightarrow{P} \theta \text{ as } n \rightarrow \infty.$$

$$\frac{\hat{\theta}_n - \theta}{\sigma_{\hat{\theta}}} \rightarrow N(0, 1),$$

where $\sigma_{\hat{\theta}}^2 = 1 / (n \mathbb{E}[-\frac{\partial^2 \log f(Y; \theta)}{\partial \theta^2}])$.

Remark: this implies that $\hat{\theta}$ is consistent and asymptotically normal. The term $n \mathbb{E}[-\frac{\partial^2 \log f(Y; \theta)}{\partial \theta^2}]$ is called **Fisher information**.

CI based on MLE

We can then obtain the following approximate large-sample $100(1 - \alpha)\%$ confidence interval for θ :

$$\hat{\theta}_n \pm z_{\alpha/2} \sqrt{1/(n\mathbb{E}[-\frac{\partial^2 \log f(Y; \theta)}{\partial \theta^2}])} \big|_{\theta=\hat{\theta}_n}.$$

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .
- We know the MLE is $\hat{p} = \frac{\sum_{i=1}^n Y_i}{n}$

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .
- We know the MLE is $\hat{p} = \frac{\sum_{i=1}^n Y_i}{n}$
- Also, $f(Y | p) = p^Y(1 - p)^{1-Y}$ for $Y = 0, 1$

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .
- We know the MLE is $\hat{p} = \frac{\sum_{i=1}^n Y_i}{n}$
- Also, $f(Y | p) = p^Y (1 - p)^{1-Y}$ for $Y = 0, 1$
- $\log f(Y | p) = Y \log p + (1 - Y) \log(1 - p)$

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .
- We know the MLE is $\hat{p} = \frac{\sum_{i=1}^n Y_i}{n}$
- Also, $f(Y | p) = p^Y (1 - p)^{1-Y}$ for $Y = 0, 1$
- $\log f(Y | p) = Y \log p + (1 - Y) \log(1 - p)$
- $\frac{\partial \log f(Y|p)}{\partial p} = \frac{Y}{p} - \frac{1-Y}{1-p}$

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .
- We know the MLE is $\hat{p} = \frac{\sum_{i=1}^n Y_i}{n}$
- Also, $f(Y | p) = p^Y (1 - p)^{1-Y}$ for $Y = 0, 1$
- $\log f(Y | p) = Y \log p + (1 - Y) \log(1 - p)$
- $\frac{\partial \log f(Y|p)}{\partial p} = \frac{Y}{p} - \frac{1-Y}{1-p}$
- $\frac{\partial^2 \log f(Y|p)}{\partial p^2} = -\frac{Y}{p^2} - \frac{1-Y}{(1-p)^2}$

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .
- We know the MLE is $\hat{p} = \frac{\sum_{i=1}^n Y_i}{n}$
- Also, $f(Y | p) = p^Y (1 - p)^{1-Y}$ for $Y = 0, 1$
- $\log f(Y | p) = Y \log p + (1 - Y) \log(1 - p)$
- $\frac{\partial \log f(Y|p)}{\partial p} = \frac{Y}{p} - \frac{1-Y}{1-p}$
- $\frac{\partial^2 \log f(Y|p)}{\partial p^2} = -\frac{Y}{p^2} - \frac{1-Y}{(1-p)^2}$
- $\mathbb{E}\left[-\frac{\partial^2 \log f(Y|p)}{\partial p^2}\right] = \mathbb{E}\left[\frac{Y}{p^2} - \frac{1-Y}{(1-p)^2}\right] = \frac{1}{p} + \frac{1}{1-p} = \frac{1}{p(1-p)}$

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .
- We know the MLE is $\hat{p} = \frac{\sum_{i=1}^n Y_i}{n}$
- Also, $f(Y | p) = p^Y (1 - p)^{1-Y}$ for $Y = 0, 1$
- $\log f(Y | p) = Y \log p + (1 - Y) \log(1 - p)$
- $\frac{\partial \log f(Y|p)}{\partial p} = \frac{Y}{p} - \frac{1-Y}{1-p}$
- $\frac{\partial^2 \log f(Y|p)}{\partial p^2} = -\frac{Y}{p^2} - \frac{1-Y}{(1-p)^2}$
- $\mathbb{E}\left[-\frac{\partial^2 \log f(Y|p)}{\partial p^2}\right] = \mathbb{E}\left[\frac{Y}{p^2} - \frac{1-Y}{(1-p)^2}\right] = \frac{1}{p} + \frac{1}{1-p} = \frac{1}{p(1-p)}$
- $\hat{\theta}_n \pm z_{\alpha/2} \sqrt{1/(n \mathbb{E}[-\frac{\partial^2 \log f(Y;\theta)}{\partial \theta^2}])} \big|_{\theta=\hat{\theta}_n}$.

Example

- For random variable with a Bernoulli distribution with parameter p . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for p .
- We know the MLE is $\hat{p} = \frac{\sum_{i=1}^n Y_i}{n}$
- Also, $f(Y | p) = p^Y (1 - p)^{1-Y}$ for $Y = 0, 1$
- $\log f(Y | p) = Y \log p + (1 - Y) \log(1 - p)$
- $\frac{\partial \log f(Y|p)}{\partial p} = \frac{Y}{p} - \frac{1-Y}{1-p}$
- $\frac{\partial^2 \log f(Y|p)}{\partial p^2} = -\frac{Y}{p^2} - \frac{1-Y}{(1-p)^2}$
- $\mathbb{E}\left[-\frac{\partial^2 \log f(Y|p)}{\partial p^2}\right] = \mathbb{E}\left[\frac{Y}{p^2} - \frac{1-Y}{(1-p)^2}\right] = \frac{1}{p} + \frac{1}{1-p} = \frac{1}{p(1-p)}$
- $\hat{\theta}_n \pm z_{\alpha/2} \sqrt{1/(n \mathbb{E}[-\frac{\partial^2 \log f(Y;\theta)}{\partial \theta^2}])} \big|_{\theta=\hat{\theta}_n}$
- The $100(1 - \alpha)\%$ CI is $\hat{p} \pm z_{\alpha/2} \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$.

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .
- We know the MLE is $\hat{\theta} = \bar{Y}$

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .
- We know the MLE is $\hat{\theta} = \bar{Y}$
- $f(Y | \theta) = \frac{1}{\theta} e^{-Y/\theta}$ for $Y = 0, 1, 2, \dots$

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .
- We know the MLE is $\hat{\theta} = \bar{Y}$
- $f(Y | \theta) = \frac{1}{\theta} e^{-Y/\theta}$ for $Y = 0, 1, 2, \dots$
- $\log f(Y | \theta) = -\log \theta - Y/\theta$

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .
- We know the MLE is $\hat{\theta} = \bar{Y}$
- $f(Y | \theta) = \frac{1}{\theta} e^{-Y/\theta}$ for $Y = 0, 1, 2, \dots$
- $\log f(Y | \theta) = -\log \theta - Y/\theta$
- $\frac{\partial \log f(Y|\theta)}{\partial \theta} = -\frac{1}{\theta} + \frac{Y}{\theta^2}$

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .
- We know the MLE is $\hat{\theta} = \bar{Y}$
- $f(Y | \theta) = \frac{1}{\theta} e^{-Y/\theta}$ for $Y = 0, 1, 2, \dots$
- $\log f(Y | \theta) = -\log \theta - Y/\theta$
- $\frac{\partial \log f(Y|\theta)}{\partial \theta} = -\frac{1}{\theta} + \frac{Y}{\theta^2}$
- $\frac{\partial^2 \log f(Y|\theta)}{\partial \theta^2} = \frac{1}{\theta^2} - 2\frac{Y}{\theta^3}$

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .
- We know the MLE is $\hat{\theta} = \bar{Y}$
- $f(Y | \theta) = \frac{1}{\theta} e^{-Y/\theta}$ for $Y = 0, 1, 2, \dots$
- $\log f(Y | \theta) = -\log \theta - Y/\theta$
- $\frac{\partial \log f(Y|\theta)}{\partial \theta} = -\frac{1}{\theta} + \frac{Y}{\theta^2}$
- $\frac{\partial^2 \log f(Y|\theta)}{\partial \theta^2} = \frac{1}{\theta^2} - 2\frac{Y}{\theta^3}$
- $\mathbb{E}\left[-\frac{\partial^2 \log f(Y|\theta)}{\partial \theta^2}\right] = \frac{1}{\theta^2}$, since $\mathbb{E}[Y] = \theta$.

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .
- We know the MLE is $\hat{\theta} = \bar{Y}$
- $f(Y | \theta) = \frac{1}{\theta} e^{-Y/\theta}$ for $Y = 0, 1, 2, \dots$
- $\log f(Y | \theta) = -\log \theta - Y/\theta$
- $\frac{\partial \log f(Y|\theta)}{\partial \theta} = -\frac{1}{\theta} + \frac{Y}{\theta^2}$
- $\frac{\partial^2 \log f(Y|\theta)}{\partial \theta^2} = \frac{1}{\theta^2} - 2\frac{Y}{\theta^3}$
- $\mathbb{E}\left[-\frac{\partial^2 \log f(Y|\theta)}{\partial \theta^2}\right] = \frac{1}{\theta^2}$, since $\mathbb{E}[Y] = \theta$.
- $\hat{\theta}_n \pm z_{\alpha/2} \sqrt{1/(n\mathbb{E}[-\frac{\partial^2 \log f(Y;\theta)}{\partial \theta^2}])} \big|_{\theta=\hat{\theta}_n}$.

Example

- For random variable with a Exponential distribution with parameter θ . If Y_1, Y_2, \dots, Y_n denote a random sample of size n from this distribution, derive a $100(1 - \alpha)\%$ confidence interval for θ .
- We know the MLE is $\hat{\theta} = \bar{Y}$
- $f(Y | \theta) = \frac{1}{\theta} e^{-Y/\theta}$ for $Y = 0, 1, 2, \dots$
- $\log f(Y | \theta) = -\log \theta - Y/\theta$
- $\frac{\partial \log f(Y|\theta)}{\partial \theta} = -\frac{1}{\theta} + \frac{Y}{\theta^2}$
- $\frac{\partial^2 \log f(Y|\theta)}{\partial \theta^2} = \frac{1}{\theta^2} - 2\frac{Y}{\theta^3}$
- $\mathbb{E}\left[-\frac{\partial^2 \log f(Y|\theta)}{\partial \theta^2}\right] = \frac{1}{\theta^2}$, since $\mathbb{E}[Y] = \theta$.
- $\hat{\theta}_n \pm z_{\alpha/2} \sqrt{1/(n\mathbb{E}[-\frac{\partial^2 \log f(Y;\theta)}{\partial \theta^2}])} \big|_{\theta=\hat{\theta}_n}$.
- The $100(1 - \alpha)\%$ CI is $\bar{Y} \pm z_{\alpha/2} \sqrt{\frac{\bar{Y}^2}{n}}$.

Summary

- General confidence intervals construction based on
 - pivotal quantities
 - MLE (likelihood function, Fisher information)

Homework

8.44, 8.102, 9.82 (b), 9.84 (a, d), 9.97 (b), 9.102